

Learning From Data

Lecture 6: Deep Neural Networks

Yang Li yangli@sz.tsinghua.edu.cn

TBSI


October 25, 2019

Today's Lecture

- ▶ Introduction to neural networks
 - ▶ Biological motivations
 - ▶ A case study
- ▶ Training a deep feedforward neural network
 - ▶ Forward pass
 - ▶ Backward propagation

Biological motivation

Schematic of a single neuron:




Each neuron takes information from other neurons, processes them, and then produces an output.

Biological motivation


How does a neuron process its input? (a *coarse* model)

- ▶ Takes the weighted average of I inputs, e.g. $z = \sum_{i=0}^I w_i(x_i)$
- ▶ Neuron fires if z is above some threshold


We call the threshold function **activation function**.



$$\text{sigmoid}(z) = \frac{1}{1+e^{-z}}$$



$$\begin{aligned}\text{tanh}(z) &= \frac{e^z - e^{-z}}{e^z + e^{-z}} \\ &= 2(\text{sigmoid}(2z)) - 1\end{aligned}$$




$$\text{ReLU}(z) = \max\{0, z\}$$

Rectifying linear unit

Biological motivation

An artificial neuron with inputs x_1, x_2 and activation function f



A single neuron is a (linear) binary classifier:

- ▶ When f is the sigmoid function, equivalent to binary softmax
- ▶ When f is the sign function, equivalent to the perceptron

Neural networks

- ▶ The goal of a neural network is to approximate some function f^* such that $y = f^*(x)$.
- ▶ The neural network defines a mapping $y = f(x; \theta)$ and learns the value of parameters θ through training.
- ▶ Define **error function** that measures prediction error of f :
e.g. a common error function used in classification is the **logarithmic loss** a.k.a. **cross-entropy loss**:

$$L = y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})$$


- ▶ $\hat{y} = f(x; \theta)$ is the predicted output
- ▶ y is the true output

A single layer of neurons are unable to approximate complex functions.

A feed forward neural network

In a **feed-forward neural network** (a.k.a. **multi-layer perceptron**), all units of one layer is connected to all of the next layer.

$$f = f^{(3)}(f^{(2)}(f^{(1)}(x)))$$



- ▶ number of layers are called **depth** of the neural network
- ▶ number of units in a layer is called **width** of a layer

The XOR problem

XOR : the exclusive or


x_1	x_2	$y = x_1 \oplus x_2$
0	0	0
0	1	1
1	0	1
1	1	0

$$f(x) = w_2^T f(W_1 x + b_1) + w_{1,0}$$

$$\text{network weights: } W_1 = \begin{bmatrix} w_{02} & w_{04} \\ w_{03} & w_{05} \end{bmatrix},$$



$$b_1 = \begin{bmatrix} w_{00} \\ w_{01} \end{bmatrix}, w_2 = \begin{bmatrix} w_{1,2} \\ w_{1,1} \end{bmatrix}$$

input layer | hidden layer | output layer



The XOR problem



Suppose f is the sign function. Here is one solution:



x_1	x_2	a_1
0	0	0
0	1	1
1	0	1
1	1	1

The XOR problem



Suppose f is the sign function. Here is one solution:



x_1	x_2	a_1	a_2
0	0	0	1
0	1	1	1
1	0	1	1
1	1	1	0

The XOR problem

Suppose f is the sign function. Here is one solution:



x_1	x_2	a_1	a_2	y
0	0	0	1	0
0	1	1	1	1
1	0	1	1	1
1	1	1	0	0


The XOR problem

Use ReLu as activation function


$$f(x) = w_2^T \max(W_1 x + b_1, 0) + w_{1,0}$$

► $W_1 = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix},$

► $b_1 = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, w_2 = \begin{bmatrix} 1 \\ -2 \end{bmatrix}$



original x space



learned activation space

Universal approximation theorem


Universal approximation theorem (Cybenko,1989; Hornik et al., 1991) A feed-forward network with a single hidden layer containing a finite number of neurons can approximate any continuous functions on compact subsets of \mathbb{R}^n , under mild assumptions on the activation function.

- ▶ First proved by George Cybenko in 1989 for sigmoid activation function;
- ▶ With one hidden layer, layer width of an *universal approximator* has to be exponentially large ← *More effective to increase the **depth** of neural networks*
- ▶ ReLU networks with width $n+1$ is sufficient to approximate any continuous function of n -dimensional input variables if depth is allowed to grow. (Lu et. al, 2017; Hanin 2018)

Overfitting

Increase the size and number of layers in a neural network,

- ▶ the **capacity** , i.e. representation power of the network increases.
- ▶ but overfitting can occur: fits the noise in the data instead of the (assumed) underlying relationship.




Regularization

One way to control overfitting in training neural networks

A common regularization approach is **parameter norm penalties**

$$\tilde{L}(w; X, y) = L(w; X, y) + \lambda \Omega(w)$$

- ▶ L2 parameter regularization: $\Omega(w) = \frac{1}{2} \|w\|_2^2 = \frac{1}{2} w^T w$ drives the weights closer to the origin



- ▶ L1 parameter regularization: $\Omega(w) = \|w\|_1 = \sum_{i=1}^k |w_i|$ drives solutions more sparse.

Forward pass and Backpropagation

See Powerpoint slides.

Practical issues

Which activation function to use?

- ▶ *sigmoid* function $\sigma(z)$: gradient $\nabla f(z)$ **saturates** when z is highly positive or highly negative. Not suitable for hidden unit activation.
- ▶ *tanh*(z): similar to identity function near 0, resembles a linear model when activation is small, performs better than sigmoid. ($\tanh(0) = 0$, $\sigma(0) = \frac{1}{2}$).
- ▶ *ReLU*(z): easy to optimize (6 times faster than sigmoid), often used with affine transformation $g(W^T x + b)$

Additional resources

Deep neural network is a relative young field with lots of empirical results.

Read more on the practical things to do for building and training neural networks:

- ▶ Stanford Class on Convolutional Neural Networks:
<http://cs231n.github.io>
- ▶ Ian Goodfellow, Yoshua Bengio and Aaron Courville, *Deep Learning*, MIT Press, 2016

Demos:

- ▶ <http://vision.stanford.edu/teaching/cs231n-demos/linear-classify/>
- ▶ <https://playground.tensorflow.org/>